

Math 158 Lecture Notes (Professor: Jacques Verstraete)

Isabelle Mills

February 24, 2024

Lecture 1: 1/9/2024

A graph is a pair (V, E) where V is a set of vertices and E is a set of unordered pairs of elements of V called edges. For $u, v \in V$, we say u and v are adjacent if $\{u, v\} \in E$.

For example: $G = (\{1, 2, 3\}, \{\{1, 2\}, \{2, 3\}\})$



A directed graph (a.k.a a digraph) is a pair (V, E) where V is a set of vertices and E is a set of ordered pairs of elements of V .

For example: $G = (\{1, 2, 3\}, \{(1, 2), (2, 3)\})$



A multigraph is a pair (V, E) where V is a set of vertices and E is a multiset of unordered pairs of elements of V .

For example: $G = (\{1, 2, 3\}, \{\{1, 2\}, \{2, 3\}, \{2, 3\}\})$



A pseudograph is like a graph and multigraph except that the pairs in E are multisets.

Essentially, an element $\{a, a\}$ can belong to E in a pseudograph. This type of edge is called a loop.

For example: $G = (\{1, 2, 3\}, \{\{1, 2\}, \{2, 3\}, \{3, 3\}\})$



If $G = (V, E)$ and $v \in V$, the neighborhood of v is $N_G(v) = \{w \in V \mid \{v, w\} \in E\}$.

The degree of v is $d_G(v) = |N_G(v)|$. Or in other words, v 's degree is equal to the number of edges connecting to v .

The Handshaking lemma states that for any graph (V, E) :

$$\sum_{v \in V} d_G(v) = 2|E|$$

The reason for this is that each edge increments the degrees of exactly two vertices. So the above sum counts every edge twice.

Lemma: Every graph has an even number of vertices with odd degrees.

Proof: We can split the vertices of any graph into two categories: those with odd degrees, and those with even degrees.

Now recall that an even number plus an even number always equals an even number, as does an odd number plus an odd number. However, an odd number plus an even number equals an odd number. Based on this fact, we can guarantee that the sum of even degrees in any graph is even. And since the sum of even degrees plus the sum of odd degrees must be even as it equals $2|E|$ by the Handshaking lemma, we thus know that the sum of odd degrees must be even. Hence, it must be the case that there are an even number of vertices with odd degree because otherwise the sum of their degrees won't be even.

A graph is called r -regular if all of its vertices have degree r .

Note that the number of edges in any n -vertex r -regular graph is $\frac{rn}{2}$.

An r -dimensional cube graph, denoted as Q_r , is a graph such that $V(Q_r)$, the set of vertices in Q_r , is equal to the set of binary strings of length r ; and $E(Q_r)$, the set of edges in Q_r , is equal to the set of pairs of binary strings which differ in only one position.



$$|V(Q_r)| = 2^r$$

$$|E(Q_r)| = \frac{2^r r}{2} = 2^{r-1} r$$

Note that Q_r is r -regular.

If $G = (V, E)$, then $H = (W, F)$ is a subgraph of G if $W \subseteq V$ and $F \subseteq E$.

If $W = V$, then H is a spanning subgraph of G (meaning that H has the same vertices as G but is lacking some of G 's edges)

We define subtracting a set of vertices from a graph as follows:

For $G = (V, E)$ and $X \subset V$, we define...

$$G - X = (V \setminus X, \{\{u, v\} \in E \mid \{u, v\} \cap X = \emptyset\})$$

We define subtracting a set of edges from a graph as follows:

For $G = (V, E)$ and $L \subset E$, we define...

$$G - L = (V, E \setminus L)$$

Lecture 2: 1/11/2024

We shall notate that H is a subgraph of G by writing $H \subseteq G$.

An induced subgraph of $G = (V, E)$ is a subgraph $G[X] = G - (V \setminus X)$ where $X \subseteq V$. Alternatively, this is called the subgraph induced by X .

Given $G = (V, E)$ and $F \subseteq E$, the subgraph spanned by F is the subgraph whose edge set is F and whose vertex set is $\bigcup_{e \in F} e$.

Here are some basic classes of graphs:

- Complete graphs / cliques, denoted K_n , are graphs where every possible edge is present between n vertices.



Note we can also draw K_4 such that there are no edge interceptions as follows:



$$|V(K_n)| = n$$

$$|E(K_n)| = \binom{n}{2} = \frac{n(n-1)}{2}$$

- A graph $G = (V, E)$ is bipartite if there exists a partition (A, B) of V such that every edge in E has one end in A and one end in B .



The partition (A, B) is called the bipartition of G . Then A and B are called the parts of G .

- A Complete bipartite graphs $K_{s,t}$, is the bipartite graph with parts A and B where $|A| = s$, $|B| = t$, and all possible edges between A and B exist.

For example, $K_{3,2}$:



- A path P_k of length k has a vertex set $V = \{v_1, v_2, \dots, v_k, v_{k+1}\}$ and an edge set $E = \{\{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{k-1}, v_k\}, \{v_k, v_{k+1}\}\}$.

Note that $|V(P_k)| = k + 1$ and $|E(P_k)| = k$.
Therefore, below would be P_3 ...



- A cycle C_k of length k has a vertex set $V = \{v_1, v_2, \dots, v_k\}$ and an edge set $E = \{\{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{k-1}, v_k\}, \{v_k, v_1\}\}$.

Note that $|V(C_k)| = k$ and $|E(C_k)| = k$.
Therefore, below would be C_4 ...



Here is some terminology before the next lemma. For the graph $G = (V, E)$...

- $\delta(G) = \min\{d_G(v) \mid v \in V\}$ is the minimum degree of G .
- $\Delta(G) = \max\{d_G(v) \mid v \in V\}$ is the maximum degree of G .
- The degree sequence of G is the sequence of degrees of vertices G in non-increasing order.

Lemma (part 1): If $G = (V, E)$ is a graph of minimum degree $k \geq 2$, then G contains a cycle of length at least $k + 1$.

Proof: Let P be a longest possible path in G , say:

$$V(P) = \{v_1, v_2, \dots, v_r\}$$

Then $N(v_r) \subseteq V(P)$. After all, if this were not the case, we'd be able to extend the path to the vertex in $N(v_r)$ but not in $V(P)$, thus contradicting the fact that P is a longest path.

Let v_i be the first neighbor of v_r along the path from v_1 to v_r . Then $\{v_i, v_{i+1}, \dots, v_r\}$ are the vertices of a cycle C .

Now note that because $N(v_r) \subseteq P$ and v_i was the first element in the path P to belong to $N(v_r)$, we know that C contains all the elements of P that $N(v_r)$ also has. So, $N(v_r) \subseteq C$.

But now note that $|N(v_r)| \geq \delta(G) = k$. Plus, v_r itself is not in $N(v_r)$. Combining these facts together, we can say that the cycle C has at least $k + 1$ vertices.

Lemma (part 2): The cycle length $k + 1$ is the longest we can guarantee based on the minimum degree of the graph being k .


Proof: Take the graph K_{k+1} which has a minimum degree k .

Obviously, the longest cycle in K_{k+1} is the cycle containing all $k + 1$ elements of K_{k+1} . Thus, we have shown that there are graphs with minimum degree k which don't have cycles of length greater than $k + 1$.

A connected graph is a graph in which any two vertices are the ends of a path.

The components of a graph are the maximal connected subgraphs. For example:

Let us define G as:



As can be seen, G has three components.

A tree is a connected graph with no cycles (a.k.a it is acyclic).

Some examples of small trees include: K_1 , K_2 , $K_{1,2}$, P_3 , and $K_{1,3}$.

Lemma: Every tree with n vertices has exactly $n - 1$ edges.

Proof: We shall proceed by induction.

If $n = 1$, the tree is K_1 , meaning that it has $0 = n - 1$ edges.

Now assume the lemma is true for all trees with n vertices, and let T be a tree with $n+1$ vertices. Then, we shall remove a vertex v of T with degree 1. (Note that we know such a vertex must exist since otherwise the minimum degree of T would be at least 2 and that would guarantee a cycle exists of at least length 3. This of course contradicts the fact that T is acyclic.)

Then $T - \{v\}$ is a tree with n vertices as it must be acyclic and connected. So by induction it has $n - 1$ edges. And because v has degree 1, we know that $|E(T)| = 1 + |E(T - \{v\})| = 1 + (n - 1) = n$.

Lemma: Any connected graph with finite vertices has a spanning tree.

Proof:

Firstly consider the case that the graph G has no cycle. Then, it is a tree by definition.

Now, consider if G has a cycle C . Then for any edge $e \in E(C)$, we have that $G - \{e\}$ is still connected. So, we can now go back to the top of the proof and ask: does $G - \{e\}$ have any cycles? We can repeatedly do this until the graph has no cycles since taking away edges does not remove any vertices.

This actually acts as an algorithm for finding a spanning tree of any connected graph.

If u and v are two vertices in a connected graph, the distance from u to v is the length of a shortest path with ends at u and v .



Let $d_G(u, v)$ be the distance between u and v .

Distance is a metric, meaning:

1. $d_G(u, v) = 0 \iff u = v$
2. $d_G(u, v) = d_G(v, u)$
3. $\forall w \in V, d_G(u, v) \leq d_G(u, w) + d_G(w, v)$

The diameter of a connected graph G is the maximum distance between any two vertices of G . Or in other words, $\max\{d_G(u, v) \mid u, v \in V(G)\}$.

The radius of G is equal to $\min\{\max\{d_G(u, v) \mid u \in V(G)\} \mid v \in V(G)\}$. What that means is that the radius of G measures the smallest distance path one could limit themselves to drawing while still being able to have that path have one end at some fixed vertex and its other end at any arbitrary vertex in the graph.

Examples:

1. The radius of K_n is 1. The diameter of K_n is n .
2. The diameter of P_k is k . The radius, can be computed as follows:
The middle vertex of a path will have the fastest access to either end of the path. So, we shall measure the radius from the vertex: $v_{\lceil \frac{k+1}{2} \rceil}$. Then, we can see that v_{k+1} is going to be a farthest element from $v_{\lceil \frac{k+1}{2} \rceil}$. So the radius of P_k equals $k + \lceil \frac{k+1}{2} \rceil$.

Now you can consider what happens when k is even and odd. But what's important is that it works out that the radius is $\lceil \frac{k}{2} \rceil$.

We can use a search tree to more generally find the radii and diameters of graphs.

Breadth-First-Search

Here's how to find a spanning tree in a connected graph with a root vertex v such that the tree "preserves" all distances from v . (This tree is called a BFS tree).

Let G be a connected graph and let $(v_1, v_2, v_3, \dots, v_n)$ be any ordering of the vertices of G .

Pick a vertex $v = v_1$ to be the root of the BFS tree.

Now, at any stage in constructing this tree, we will have a vertex set $V(T) = \{v_1, v_2, \dots, v_k\}$ (when we first start, $V(T)$ will only contain v_0 . So don't worry about that). Now if $V(T) = V(G)$ we can stop. Otherwise though, we can say that there is a smallest integer i such that for $v_i \in V(T)$, $N(v_i) \setminus V(T) \neq \emptyset$. Choose v_{k+1} to be the smallest neighbor (by the ordering of $V(G)$) of v_i not in T and add the edge $\{v_i, v_{k+1}\}$ to T . Then we repeat this paragraph.

Beware the ordering we are creating in our tree will often be different from the order of the graph you started with.

Worked Demonstration:



Properties of BFS:

- If the root is v , then $d_T(v, w) = d_G(v, w)$. In other words, a BFS tree preserves distances from its root.
- The Tree with root v has layers $N_i(v) = \{w \in V(G) \mid d_G(v, w) = i\}$. Furthermore all edges in the original graph either stay inside a single layer $N_i(v)$ or go between adjacent layers (i.e. from $N_i(v)$ to $N_{i+1}(v)$).
If an edge did "jump over" a layer, that violate the fact that distance is a metric.
- The diameter of G equals the maximum number of layers of all BFS trees (not including the 0-layer).
- The radius of G equals the minimum number of layers of all BFS trees (also not including the 0-layer).

Lecture 3: 1/16/2024

Note that a tree is "minimally connecting" as subtracting any edge from a tree will produce a disconnected graph.

We know this is the case because if we could remove an edge and still have the graph be connected, then that would imply the existence of a path between two neighboring vertices that doesn't go through their shared edge. But then, we'd be able to make a cycle subgraph by adding their shared edge to that path.

Depth-First-Search

Here is alternate algorithm for generating a spanning tree of a connected graph. A resulting tree of this algorithm is called a DFS tree.


Let G be a connected graph and let (v_1, v_2, \dots, v_n) be any ordering of the vertices of G .

Pick a vertex $v = v_1$ to be the root of the DFS-tree.

Now, at any stage in constructing this tree, we will have a vertex set $V(T) = \{v_1, v_2, \dots, v_k\}$. If $V(T) = V(G)$, we can stop. Otherwise though, we select i to be the largest integer such that for $v_i \in V(T)$, $N(v_i) \setminus V(T) \neq \emptyset$. Then, choose v_{k+1} to be the smallest neighbor (by the ordering of $V(G)$) of v_i not in $V(T)$ and add the edge $\{v_i, v_{k+1}\}$ to T . Then we repeat this paragraph.

Once again beware the ordering we are creating in our tree will typically be different from the order of the graph you started with.

Worked Demonstration:



Theorem: A graph is bipartite if and only if it contains no odd cycles.

Proof:

(\implies) First note that an odd cycle isn't bipartite. Thus, any graph containing an odd cycle is not bipartite.

(\impliedby) Now supposed we are given some graph G with no odd cycles.

Then, assuming G is connected (if G isn't connected, we can break G up into its component subgraphs and do this process for each component), we can construct a BFS-tree in G rooted at some $v \in V(G)$. Let us name this tree T .

Now as noted before, T will have layers L_i where each $L_i = \{u \in V(G) \mid d_G(v, u) = i\}$. Using those layers, we can partition T into two subsets A and B where A is the union of all L_i where i is even and B is the union of all L_j where j is odd. So, T is clearly bipartite.

Now, let's reinsert the removed edges from G back into T . Note that for each re-inserted edge e , it must be the case that either e is a subset of some L_i or that e goes between some L_i and L_{i+1} . Importantly, edges of the latter case do not violate our partition. So, if all the edges in $E(G) \setminus E(T)$ go between layers, then we can conclude that G is definitely bipartite just like T .

With that, we now intend to show that an edge G having an edge belonging to a single layer L_i guarantees that G contains an odd cycle.

Assume the graph G has an edge $\{u, w\} \subseteq L_i$ where L_i is the i th layer of a BFS tree rooted at v . Then, we know that there exists a path P_1 contained in that BFS tree going from v to u and a path P_2 contained in that BFS tree going from v to w . In order to draw a cycle from this information, let x be the vertex of some L_j such that $x \in V(P_1)$, $x \in V(P_2)$, and j is as large as possible. That way, by defining the subpaths P_1' going from x to u and P_2' going from x to w , we can get the following cyclic subgraph of G :

$$C = (V(P_1') \cup V(P_2'), E(P_1') \cup E(P_2') \cup \{u, w\})$$

However, now note that $|E(P_1')| = |E(P_2')| = i - j$. Hence, $|E(C)| = 2(i - j) + 1$, which in turn means that C has an odd number of edges. So, we have shown that if a graph G contains an edge within a single layer L_i , then we can give an example of an odd cycle within G .

So in conclusion, if we assume G has no odd cycles, then G can't have any edges which are subsets of a single layer L_i . But that means that every edge in G respects the partition we made to show that T is bipartite. So, G must also be bipartite with the same partition as T . ■

A Hamiltonian cycle is a spanning cycle of a graph. We say a graph is Hamiltonian if it contains such a cycle.

A Hamiltonian path is a spanning path of a graph. We say a graph is traceable if it has a hamiltonian path.

A walk is a sequence of vertices and edges: i.e. $(v_1, \{v_1, v_2\}, v_2, \{v_2, v_3\}, \dots)$

Note that a walk can go over the same edge or vertex multiple times.

A trail is a walk with no repeated edge.

Interestingly, all paths are trails and all trails are walks. So a trail is kind of a middle concept between being a walk or a path.

A tour is a trail with the same first and last vertex.

So, all cycles are tours and all tours are walks.

An Eulerian tour of a graph is a tour which contains all the edges of the graph.

For context, the name Eulerian is in honor of Leonhard Euler because he was the first mathematician to ask when a graph would have an Eulerian tour (look up the Seven Bridges of Königsburg problem).

We call a graph Eulerian if for every vertex v in the graph: $d(v)$ is even.

Lecture 4: 1/18/2024

Before finding conditions for the existence of an Eulerian tour of a graph, let's establish some terminology for digraphs so that we can study Eulerian tours in digraphs as well.

Firstly, a walk, trail, and tour are defined almost identically in a digraph as in a graph. The one difference is that given some edge (u, v) of a digraph, a walk, trail, and path are only allowed to traverse that edge going from u to v .

Given a digraph (V, E) , the "out" and "in" neighborhoods of a vertex $v \in V$ are:

out: $N^+(v) = \{w \in V \mid (v, w) \in E\}$

in: $N^-(v) = \{w \in V \mid (w, v) \in E\}$

Similarly, the out-degree and in-degree of $v \in V$ are:

out: $d^+(v) = |N^+(v)|$

in: $d^-(v) = |N^-(v)|$

A digraph is called Eulerian if for each vertex v , $d^+(v) = d^-(v)$.



$$\begin{array}{lll} d^+(v_1) = 1 & d^+(v_2) = 0 & d^+(v_3) = 2 \\ d^-(v_1) = 1 & d^-(v_2) = 2 & d^-(v_3) = 0 \end{array}$$

This digraph is not Eulerian.

An orientation of a graph G is a digraph with the same vertices as G but where each $\{u, v\} \in E(G)$ is replaced with either (u, v) or (v, u) .

The underlying graph (or multigraph) of a digraph G is the graph (or multigraph) such that $\{u, v\}$ is an edge whenever $(u, v) \in G$.

Theorem:

1. A graph has an Eulerian tour if and only if it is connected and Eulerian.
2. A digraph has an Eulerian tour if and only if it has a connected underlying graph and if it is Eulerian.

Proof of statement 1:

(\implies) If G has an Eulerian tour $v_0e_0v_1e_1v_2e_2\ldots v_ke_kv_0$, then for any v_i , the tour has to use an edge into v_i and an edge out of v_i each time it visits v_i . So, $d(v_i)$ is even for all i .

(\impliedby) Now suppose G is connected and all vertices have even degree. Then let $T = v_0e_0v_1e_1v_2e_2\ldots v_le_lv_{l+1}$ be a longest trail in G .

If $v_{l+1} \neq v_0$, then we know that v_{l+1} has an odd degree in T as the trail goes into v_{l+1} and doesn't leave. However, because we assumed that all vertices in G have an even degree, we know there must be an even number of edges coming out of v_{l+1} . So, we can add another edge to our trail to get a longer trail. But this contradicts our assumption that T is a longest trail in G . Hence, we conclude that $v_{l+1} = v_0$, meaning T is a tour.

Now consider if T is not an Eulerian tour. In that case, there is an edge of G not in T . Additionally, because G is connected, we know that that edge will have the form $e = \{v_i, w\}$. So now consider a new trail: T' defined as $v_ie_iv_{i+1}e_{i+1}\ldots v_0e_0v_1e_1\ldots v_iw$. Importantly, T' is a longer trail than T . So we have a contradiction as T is not a longest trail.

Therefore, the longest trail T in G must be an Eulerian tour.

The proof of statement 2 is nearly identical.

Note that this proof can be interpreted as giving an algorithm for finding an Eulerian tour.

1. Make a trail T .
2. Add edges to T until you get stuck at a vertex. Then you know that your trail forms a tour.
3. If T is not an Eulerian tour, then going by the steps in the proof above, define T' . Then do step 2. on T' .
4. If T is an Eulerian tour, you're done.

A harder problem is whether a graph has a Hamiltonian (spanning) cycle or not.



Dirac's Theorem: Let $n \geq 3$ and let G be an n -vertex graph of minimum degree at least $\frac{n}{2}$. Then G is hamiltonian.

Proof:

Suppose the theorem is false. Let G be a counter-example with as many edges as possible (a maximal counter-example). Then, we know there exists an edge $\{u, v\} \notin E(G)$ as G cannot equal K_n since K_n is hamiltonian. Furthermore, we know that $G + \{u, v\}$ is hamiltonian since G was maximal. So, there is a hamiltonian cycle $C \subseteq G + \{u, v\}$ using the edge $\{u, v\}$. This in turn means that $C - \{u, v\}$ is a hamiltonian path belonging to G .

Let $P = v_1 e_1 v_2 e_2 \dots e_n v_n$ be the hamiltonian path in G from u to v and let $N^+(w)$ denote the set of vertices immediately following a neighbor of w on the path P . In other words: $N^+(w) = \{v_{i+1} \mid v_i \in N_G(w)\}$.

By the theorem's assumption about the minimum degree of the graph, we know that $|N_G(u)| \geq \frac{n}{2}$. Meanwhile on the other end of P , since $v = v_n \notin N_G(v)$, we know that every neighbor of v has an element following it on the path. So $|N^+(v)| \geq \frac{n}{2}$. Thus $|N_G(u)| + |N^+(v)| \geq n$. But now note that u does not belong to either of the above sets. So $|N_G(u) \cup N^+(v)| \leq n - 1$. As a consequence of this, $|N_G(u) \cap N^+(v)| > 0$.

Let v_i be a vertex belonging to $(N_G(v) \cap N^+(u))$. Then we can draw a cycle in G visiting all the vertices of G in the following order:

$$v_1, v_2, \dots, v_{i-1}, v_n, v_{n-1}, \dots, v_i, v_1$$

This contradicts our assumption that G would be a counter example and thus not Hamiltonian. So, we assume no such counter example exists.

We can also show that Dirac's Theorem gives the best possible minimum degree for a graph to be guarenteably Hamiltonian. Consider a graph G containing two copies of K_m sharing exactly one vertex. In that case, $n = |V(G)| = 2m - 1$ and $\delta(G) = m - 1 = \frac{n-1}{2}$. However, this graph does not have a spanning subcycle as any spanning cycle would have to cross that shared vertex twice.

Let P be a longest path in G going from a vertex u to a vertex v . Additionally, for any $w \in V(P)$, let w^+ be the vertex following w as one travels from u to v along P . Importantly, since P is a longest path, we know that $(N_G(u) \cup N_G(v)) \subseteq V(P)$. Furthermore, for each $w \in N_G(v)$, we can define $Q = P - \{w, w^+\} + \{v, w\}$ where Q is a longest path of G going from u to w^+ instead of u to v . We call Q a rotation of P at v .

Pósa's Rotation Lemma: Suppose G is a graph and for every $S \subseteq V(G)$ with $|S| \leq t$, $|N(S)| > 2|S|$. Then G contains a path of length $3t + 1$.

$N(S)$ is referring to the union of the neighborhoods of each $v \in S$ minus any vertices in S .

Proof:

Let P be a longest path ending at a vertex v . Also, let S be the set of end vertices of all possible longest paths that could be obtained through any number of rotations starting with P . Finally, let S^+ and S^- denoted the vertices of P immediately after and immediately before vertices in S respectively.

Obviously, $|S^+| \leq |S|$ and $|S^-| \leq |S|$ as all vertices except the first and last vertex of P have exactly one vertex before and after them in P . Also note that $N(S) \subseteq S^+ \cup S^-$. This is because if there did exist $w \in N(S)$ such that $w \notin S^+ \cup S^-$, then we would know that no rotation of P made it so that w was not proceeded by w^- and followed by w^+ . So, doing a rotation with the vertex w , we would show that either w^+ or w^- belonged to S , thus reaching a contradiction.

Overall, this means that $|N(S)| \leq |S^+ \cup S^-| \leq |S^+| + |S^-| \leq 2|S|$. However, note that by the theorem's assumption about G , we know that $|S| \geq t$ because otherwise we'd have that $2|S| < |N(S)|$. So, let T be a subset of S such that $|T| = t$. Then, because T and $N(T)$ are disjoint subsets of $(S \cup N(S))$ which itself is a subset of $V(P)$, we know that $|V(P)| \geq |T| + |N(T)|$. And since $|N(T)| > 2|T| = 2t$ by the theorem's assumption about G , we thus can say that $|V(P)| > 3t$. ■

Lecture 5: 1/23/2024

Theorem: If for every set S of vertices in a graph G , we have that $|N(S)| \geq \min \{2|S| + 1, |V(G) \setminus S|\}$, then G has a hamiltonian path.

Proof:

Once again let us define P as a longest path of G , as well as S as the set of end vertices of all possible rotations of P . Then by the same reasoning as before, we know that $|N(S)| \leq 2|S|$. Therefore, since $|N(S)|$ isn't greater than or equal to $2|S| + 1$, we know by the assumption of the theorem that $N(S)$ is greater than or equal to $|V(G) \setminus S|$.

Now $S \cup (V(G) \setminus S) = V(G)$ and $(S \cup N(S)) \subseteq V(P)$. Additionally, S and $(V(G) \setminus S)$ are disjoint to each other, as is S and $N(S)$. Thus, we can say that

$$|V(G)| = |S| + |(V(G) \setminus S)| \geq |S| + |N(S)| \leq |V(P)|$$

Therefore the longest path P must cover every vertex of G , meaning that it is a Hamiltonian path. ■

This is what is used to find Hamiltonian paths in random graphs.

A graph is uniquely Hamiltonian if it has exactly one Hamiltonian cycle.

Theorem: If all vertices in a graph G have odd degree, then every edge is in an even number of hamiltonian cycles.

In other words such a graph is not uniquely Hamiltonian.

Proof:

If the graph G in the theorem has no hamiltonian cycles, then we're done. Every edge is in 0 hamiltonian cycles.

Now pick an edge and suppose that there is a Hamiltonian cycle C containing it. We'll call that edge $\{u, v\}$ and let w be the vertex coming before u on C .

Then, let us define a new graph H whose vertices are Hamiltonian paths in G which start with the edge $\{u, v\}$. For example, $(C - \{u, w\}) \in V(H)$. Additionally, let $\{P, Q\}$ be an edge of H if P and Q are rotations of each other.

If $P \in V(H)$ is a hamiltonian path in G ending in a vertex $x \in V(G)$, then:

$$d_H(P) = \begin{cases} d_G(x) - 1 & x \notin N_G(u) \\ d_G(x) - 2 & x \in N_G(u) \end{cases}$$

Essentially, for every edge connecting to x except for the one already used by P , there is a rotation of P . We can then say that that rotation is a vertex of H if it includes the edge $\{u, v\}$ (In other words, a rotation including the edge $\{u, x\}$ would not be included in H).

If x is adjacent to u , then $P + \{\{u, x\}\}$ is a hamiltonian cycle containing $\{u, v\}$.

Now here is the clever part: since $d_G(x)$ is assumed to be odd for all $v \in V(G)$, we have that $d_H(P)$ is even if x is not adjacent to u and odd if x is adjacent to u . But now note that every graph has an even number of vertices of odd degree because of the handshaking lemma. So, there must be an even number of paths including the edge $\{u, v\}$ and ending in a vertex x such that x is adjacent to u . Or in other words, G has an even number of hamiltonian cycles containing the edge $\{u, v\}$. ■

For example, by the above theorem we know that this graph is not uniquely Hamiltonian.



One note about the above theorem is that it can be interpreted as giving an algorithm for finding a second hamiltonian cycle given one hamiltonian cycle in a graph where all degrees are odd.

A matching in a graph is a set of vertex disjoint edges in the graph.

A vertex is saturated by a matching if one of the edges of the matching contains the vertex. Otherwise, we say the vertex is exposed by the matching.



The matching shown to the left is:
 $M = \{\{1, 2\}, \{3, 4\}, \{5, 6\}\}$

In the matching to the left, 7 is exposed and all other vertices are saturated.

A maximum matching in a graph is a matching with the maximum number of edges.

A perfect matching (or 1-factor) is a matching which saturates all the vertices of a graph.

Proposition: for a graph to have a perfect matching, it must have an even number of vertices.

Lecture 6: 1/30/2024

Hall's Theorem: Let G be a bipartite graph with parts A and B . Then G has a matching saturating A if and only if for every set $X \subseteq A$,

$$|N_G(X)| \geq |X|.$$

Note: we refer to the below statement as Hall's condition:
 $\forall X \subseteq A, |N_G(X)| \geq |X|.$

Proof:

(\implies) If G has a matching M containing A , then for any $X \subseteq A$, we trivially have that $|N_G(X)| \geq |N_m(X)| = |X|$.

(\impliedby) To prove the other way, we proceed by induction on A while assuming Hall's condition is true.

Base case: assume that $|A| = 1$. In that case, because Hall's condition is assumed to be true, we know that $|N_G(A)| \geq |A| = 1$. Thus, for some $a \in A$ and $b \in B$, we know there exists an edge $\{a, b\} \in E(G)$. The matching of just that edge saturates A .

Induction step: Assume that Hall's theorem works if $1 \leq |A| < n$. Then assume that $|A| = n$. Since we are assuming Hall's condition to be true, we know that for all proper subsets $X \subset A$, either $|N_G(X)| > |X|$ or $|N_G(X)| = |X|$ is true.

- **Case 1:** For all proper subsets X of A , $|N_G(X)| > |X|$...
Pick an edge $\{a, b\} \in E(G)$ and consider $H = G - \{a\} - \{b\}$. We know that H will be a bipartite graph with two sets of vertices: $A' = A \setminus \{a\}$, and $B' = B \setminus \{b\}$. Additionally, given any $X \subseteq A'$, we know that $|N_H(X)| \geq |N_G(X)| - 1 \geq |X|$. So, by our inductive hypothesis, H has a matching covering A' . Now add the edge $\{a, b\}$ to this matching to get a matching in G covering A .
- **Case 2:** There exists a proper subset X of A such that $|N_G(X)| = |X|$...
In this second case our reasoning for case 1 breaks down because given $X \subseteq A'$, we can no longer guarantee that $|N_G(X)| - 1 \geq |X|$. So, using that same set X , consider G_1 the induced graph of $N_G(X) \cup X$ and G_2 equal to $G - V(G_1)$.

For any $Y \subseteq X$, we have that $|N_{G_1}(Y)| = |N_G(Y)| \geq |Y|$ due to Hall's condition. Thus, by our inductive hypothesis there exists a matching M_1 saturating X in G_1 .

Additionally, for any $Y \subseteq A \setminus X$, consider $N_G(Y \cup X)$. By assuming Hall's condition, we know that $|N_G(Y \cup X)| \geq |Y \cup X|$. And, because X and Y are disjoint, $|Y \cup X| = |Y| + |X|$. On the other hand, we also know that $N_G(Y \cup X) = N_{G_2}(Y) \cup N_G(X)$. So, $|N_{G_2}(Y)| + |N_G(X)| \geq |Y| + |X|$. Finally, because $|N_G(X)| = |X|$, we can cancel terms to get that: $|N_{G_2}(Y)| \geq |Y|$. Hence by our inductive hypothesis, we know there exists a matching M_2 saturating $A \setminus X$ in G_2 .

We now can combine M_1 and M_2 in order to get a matching in G covering A .

Side proposition: if in addition to Hall's condition holding, $|A| = |B|$, then the matching for G generated by the above proof is perfect. Meanwhile, if $|A| \neq |B|$, then it is impossible to make a perfect matching. To prove this, assume without loss of generality that $|A| < |B|$. Then, because $|N_G(B)|$ is at most $|A|$, we know that $|N_G(B)| < |B|$. So no matching can exist covering B .

One application of Hall's Theorem is the "systems of distinct representatives" problem.

Let S_1, S_2, \dots, S_n be committees and let $s_i \in S_j$ refer to some person in the j th committee. Then, is it possible to select a representative s_i in each committee S_j so that each s_i is distinct?

To answer this, note that we can represent the above situation as a bipartite graph where vertex set A contains people and vertex set B contains committees. Then for a person $a \in A$ and committee $b \in B$, we add the edge $\{a, b\}$ if person a is in committee b .



Theorem: Committees S_1, S_2, \dots, S_n have a system of distinct representatives if and only if for all $X \subseteq \{1, 2, \dots, n\}$:

$$\left| \bigcup_{i \in X} S_i \right| \geq |X|$$

A 1-factorization of a graph is a collection of pairwise edge-disjoint 1-factors M_1, M_2, \dots, M_r such that $E(G) = M_1 \cup M_2 \cup \dots \cup M_r$.

Theorem: Every r -regular bipartite graph has a 1-factorization if $r \geq 1$.

Proof:

Let $X \subseteq A$. Then the number of edges going out from X to B , written as $e(X, B)$, is equal to $r |X|$. Meanwhile, the number of edges going from $N(X)$ back to A is equal to $e(N(X), A) = r |N(X)|$.

Now note that every edge counted in $e(X, B)$ is also counted in $e(N(X), A)$. Therefore, $r |X| = e(X, B) \leq e(N(X), A) = r |N(X)|$. And finally as $r > 0$, we have that $|X| \leq |N(X)|$. Thus Hall's condition is satisfied.

Because Hall's condition is satisfied, we know by Hall's theorem that there is a matching M_1 saturating A . Additionally, because the total number of edges in the graph is equal to $e(A, B) = r |A| = r |B|$, we know that $|A| = |B|$. So the matching we found that saturates A also saturates B .

Now, that we've found a 1-factor M_1 , let's now subtract M_1 from our bipartite graph. Then, the resulting subgraph is $(r - 1)$ regular. So, we can apply the above reasoning again and again to get more 1-factors. And importantly, all of these 1-factors must be disjoint.

After r iterations, our bipartite graph will have no edges left and we will have created a collection of r many 1-factors whose union is the original edge set. ■

Lecture 7: 2/1/2024

For a graph G , let $\text{odd}(G)$ denote the number of components of G with an odd number of vertices.

If G has a perfect matching, then $\text{odd}(G) = 0$. Furthermore, for any set $S \subseteq V(G)$:

$$\text{odd}(G - S) \leq |S|$$

To understand why, let ϵ be the minimum number of exposed vertices in $G - S$ and consider that:

- We know $\epsilon \leq |S|$ because if M is a perfect matching of G , then the number of exposed edges by the matching: $M \cap E(G - S)$ is at most equal to $|S|$.
- Meanwhile, because no odd component of $G - S$ can have a perfect matching, we know that each odd component of $G - S$ must contribute at least one exposed vertex to an optimal matching of $G - S$. So $\text{odd}(G - S) \leq \epsilon$.

Combining the two bounds, we have $\text{odd}(G - S) \leq \epsilon \leq |S|$. Therefore, $\text{odd}(G - S) \leq |S|$.

I do not understand how the professor saw this as an obviously apparent truth that didn't need a proof but whatever.

This condition is called Tutte's condition. As shown above, it is a necessary condition for a graph to have a perfect matching. However, it happens to also be a sufficient condition.

Theorem: A graph G has a perfect matching if and only if $\text{odd}(G - S) \leq |S|$ for every set $S \subseteq V(G)$.

Proof:

We already showed that Tutte's condition was necessary for a graph to have a perfect matching. So, what we need to show now is that assuming Tutte's condition holds, you can construct a perfect matching.

We shall proceed by induction. Firstly note that if $S = \emptyset$, then there must be zero odd components. Therefore, Tutte's condition already limits us to focusing on graphs with an even total number of vertices. So, let our base case be when $|V(G)| = 2$. In that case $G = K_2$ and so G has a perfect matching.

Now suppose $V(G) > 2$ and that our theorem holds for any graph with less than $V(G)$ vertices.

Let S be the largest subset of $V(G)$ such that $|S| = \text{odd}(G - S)$. We know that S is nonempty because if $|S|$ equals 1, the equality must hold.

$|V(G)|$ having to be even implies that $|V(G)| - 1$ is odd. So $G - S$ must have an odd number of odd component. But by Tutte's condition, we know the number of odd components is at most 1. Hence, equality must hold when $|S| = 1$.

This guarantees that $V(G - S) \leq V(G)$.

Let's now consider any even component of $G - S$ which we shall denote H . We claim that H has a perfect matching.

Consider any $R \subseteq V(H)$. Because R only contains vertices from an even component of $G - S$, we have that every odd component in $G - S$ is also in $G - (R \cup S)$. This combined with Tutte's condition tells us that:

$$\text{odd}(H - R) + \text{odd}(G - S) = \text{odd}(G - (R \cup S)) \leq |S| + |R|$$

However, since we specified that $|S| = \text{odd}(G - S)$, we can cancel out terms to get that $\text{odd}(H - R) \leq |R|$. So Tutte's condition holds for H . And since $V(H) < V(G)$, we can thus conclude by our inductive hypothesis that H has a perfect matching.

Meanwhile, consider any odd component of $G - S$ which we shall denote F . We claim that for any $v \in V(F)$, we know that $F' = F - \{v\}$ has a perfect matching.

Because $V(F') < V(G)$, we know by our inductive hypothesis that if F' does not have a perfect matching, then Tutte's condition must not hold. Or in other words, there exists $Q \subseteq V(F')$ such that $\text{odd}(F' - Q) > |Q|$.

Now note that for any $R \subseteq V(F)$, because $|V(F)|$ is odd ($\equiv 1 \pmod{2}$):

- $|R| \stackrel{2}{\equiv} 0 \implies |V(F - R)| \stackrel{2}{\equiv} 1 \implies \text{odd}(F - R) \stackrel{2}{\equiv} 1$
- $|R| \stackrel{2}{\equiv} 1 \implies |V(F - R)| \stackrel{2}{\equiv} 0 \implies \text{odd}(F - R) \stackrel{2}{\equiv} 0$

So $\text{odd}(F - R) + |R| \stackrel{2}{\equiv} |V(F)| \stackrel{2}{\equiv} 1$.

This means that $\text{odd}(F - (Q \cup \{v\})) + |Q| + 1 \stackrel{2}{\equiv} 1$, which in turn says that $\text{odd}(F' - Q) + |Q| \stackrel{2}{\equiv} 0$. So both $\text{odd}(F' - Q)$ and $|Q|$ must have the same parity, meaning that $\text{odd}(F' - Q) > |Q| \implies \text{odd}(F' - Q) \geq |Q| + 2$.

Also observe that $\text{odd}(G - S \cup \{v\} \cup Q) = \text{odd}(G - S) - 1 + \text{odd}(F' - Q)$. This is because every odd component of $G - S$ except for F is also an odd component of $G - S \cup \{v\} \cup Q$.

Therefore we can say by Tutte's condition that:

$$|S| + |Q| + 1 \geq \text{odd}(G - S) - 1 + \text{odd}(F' - Q) \geq \text{odd}(G - S) + |Q| + 1.$$

Therefore we've shown that $\text{odd}(G - S \cup \{v\} \cup Q) = |S \cup \{v\} \cup Q|$. However, this contradicts our previous assertion that S is the largest subset of $V(G)$ such that $\text{odd}(G - S) = |S|$. So we have shown that Q cannot exist, meaning that F' has a perfect matching.

Finally, let C be the set of vertices consisting of one vertex from each odd component of $G - S$ such that each vertex in C has an edge connecting to S . Then consider the bipartite graph $G(S, C)$ consisting of every edge in G going between C and S . We claim that this bipartite graph has a perfect matching.

For every $X \subseteq C$, we know that removing $N(X)$ will make the components which each $x \in X = C$ belong to odd. So $|X| = \text{odd}(G - N(X))$. And by Tutte's condition: $\text{odd}(G - N(X)) \leq |N(X)|$. So $|X| \leq |N(X)|$, meaning that by Hall's theorem, $G(S, C)$ has a perfect matching.

Combining all the perfect matchings we made for $G(S, C)$ and each H and F' , we get a perfect matching for all G . ■

A bridge is an edge which when removed increases the number of components of a graph. Meanwhile, a cubic graph is a graph that is 3-regular.

Petersen's Theorem: Every bridgeless cubic graph G has a perfect matching.

We check that for all $S \subseteq V(G)$, we have that $\text{odd}(G - S) \leq |S|$.

By handshaking lemma, each component of G must have an even number of vertices. Thus, if $S = \emptyset$, then $\text{odd}(G - S) = \text{odd}(G) = 0$.

Now suppose if $S \neq \emptyset$. Then each component of $G - S$ sends at least two edges to S since G is bridgeless. However, consider a hypothetical odd component H of $G - S$. Because the sum of degrees of H must be even, we know that H must send an odd number of edges to S . Therefore if every component of $G - S$ was odd, then the number of edges coming into S must be at least $3(\text{odd}(G - S))$. But, because G is 3-regular, the max number of edges that S could possibly accept is $3|S|$. So, $\text{odd}(G - S) \leq |S|$.

Matching Algorithms:

Let G be a graph and M a matching in G .

- An M -alternating path is a path whose edges are alternately in and not in M .
- An M -augmenting path is an alternating path whose first and last edge are not in M .



Theorem: If G is a graph and M is a matching in G , then M is a maximum matching if and only there is no augmenting path between two vertices not covered by M .

Proof:


(\implies) the contrapositive of this is easy to show. If such an augmenting path does exist, then we can draw a larger matching by doing this:



(\impliedby) we're not rigorously showing this direction for some reason in this class...

Edge Coloring:

A proper edge coloring of a graph G is a map $c : E(G) \longrightarrow S$, where S is a set such that $c(e) \neq c(f)$ whenever $e \cap f \neq \emptyset$.



This is an edge coloring of C_7 .
 $S = \{\text{orange}, \text{red}, \text{pink}\}$ and c maps each edge of C_7 to an element of S as shown to the side.

The edge-chromatic number $\chi'(G)$ is the minimum k for which G has a proper k -edge-coloring.

Observe that $\chi'(G) \geq \Delta(G)$ (the maximum degree of G). After all, each edge coming off the vertex with the max degree have to have a different color.

König's Theorem: Let G be a bipartite multigraph. Then $\chi'(G) = \Delta(G)$.

Proof:

By taking two copies of G and adding multiple edges between copies of the same vertex, we can get a $\Delta(G)$ -regular-multigraph. Then, by the theorem on page 20 (which I realize we only explicitly proved for normal graphs but all the logic we used should also work for multigraphs...), we know that this new graph has $\Delta(G)$ disjoint perfect matches. Then assign colors to the original edges of G based on which perfect matching that edge wound up in.

A different proof is in the textbook...

Lecture 8: 2/6/2024

We define $\mu(G)$ to be the size of a maximum matching of a graph G . Meanwhile we define $\text{ex}(G)$ to be the number of vertices exposed by a maximum matching of G .

König-Ore Theorem: Let G be a bipartite graph with parts A and B . Additionally, define $\text{ex}(G, A) = |A| - \mu(G)$. In other words, this is the number of vertices in A exposed by a maximum matching. Then:

$$\text{ex}(G, A) = \max_{S \subseteq A} \{|S| - |N(S)|\}$$

Proof:

Denote the right hand side of the above equation by d . Then consider adding d vertices to B which are all adjacent to all vertices of A . Hall's condition is trivially true in this new graph, meaning that this new graph has a matching covering all vertices of A . Therefore, when we remove those added vertices and edges from our matching over A , we will still have a matching covering at least $|A| - d$ vertices in A . So, $\text{ex}(G, A) \leq d$.

On the other hand, if M is a matching of size: $|A| - \text{ex}(G, A)$, meaning that M is a maximum matching, then each $S \subseteq A$ has at least $|S| - \text{ex}(G, A)$ neighbors in B . Or in other words, $|N(S)| \geq |S| - \text{ex}(G, A)$. We know this because Hall's condition applies to A if we remove the exposed vertices in A . Then, by subtracting $\text{ex}(G, A)$ from $|S|$, we are effectively canceling the influence of those exposed points which cause Hall's Theorem to fail. We can rewrite the above formula as $\text{ex}(G, A) \geq |S| - |N(S)|$. And since this is true for all S , we thus know that $d = \max_{S \subseteq A} \{|S| - |N(S)|\} \leq \text{ex}(G, A)$.

So in conclusion: $d \leq \text{ex}(G, A) \leq d \implies d = \text{ex}(G, A)$.

Here are two other theorems given without proof:

Tutte-Berge Formula: For any multigraph G :

$$\text{ex}(G) = \max_{S \subseteq V(G)} \{\text{odd}(G - S) - |S|\}$$

Vizing's Theorem: For every graph G of maximum degree Δ , either $\chi'(G) = \Delta$ or $\chi'(G) = \Delta + 1$.

if $\chi'(G) = \Delta$, then G is referred to as class 1. Meanwhile,
if $\chi'(G) = \Delta + 1$, then G is referred to as class 2.

That said, it is an NP-complete problem to generally determine if a graph is class 1 or not.


An embedding of a graph $G = (V, E)$ is a function $f : V \cup E \rightarrow \mathbb{R}^2 \cap \mathcal{C}$ (where \mathcal{C} is the set of continuous curves in \mathbb{R}^2) such that f is injective, $f(v)$ is a point in \mathbb{R}^2 for all $v \in V$, and $f(\{u, v\})$ is a continuous curve in \mathbb{R}^2 with ends at $f(u)$ and $f(v)$ for all $\{u, v\} \in E$.

A graph is a planar if we can choose f such that for all distinct $e, e' \in E$, $f(e)$ only intersects $f(e')$ at its end points if at all.

In other words, G is planar if we can draw it without crossings (and we won't get into the topology of what that means...)

A plane graph is a graph which we have embedded in the plane without crossings.

A subdivision of a graph H is a graph obtained by replacing every edge $e = \{u, v\}$ with a path P_e which ends at u and v such that $V(P_e) \cap V(P_f) = e \cap f$ for all $e, f \in E(H)$.



Kuratowski's Theorem: A graph is planar if and only if it does not contain a subdivision of K_5 or $K_{3,3}$.

The proof for this is beyond the scope of this class.

Let G be a plane graph. Then the faces of G are the maximal connected regions of $\mathbb{R}^2 \setminus G$. We refer to the set of faces of G as $F(G)$.

The plane graph below has four faces:




Beware that the same graph can have multiple plane drawings.

If G is connected, then the boundary walk of a face of G is the shortest closed walk which uses every edge on the topological boundary of any face.

The degree of a face is the length of its boundary walk.

The middle face the graph below has degree 10...



The two vertices protruding inside the face have to be traversed twice by any walk going around the boundary of the internal face.

Proposition: The degree of the only face in a plane graph of a tree is twice the tree's number of edges.

Handshaking Lemma: $\sum_{f \in F(G)} \text{degrees}(f) = 2|E(G)|$

Intution why:

Each edge has a face on its left hand and right hand side (it might be the same face on both side). Thus the boundary walk of the face on the edge's left hand side will have to traverse that edge once, whereas the face on the edge's right hand side will also have to traverse that edge once.

Euler's formula: Let G be a connected plane graph. Then:

$$|V(G)| - |E(G)| + |F(G)| = 2$$

Proof: (we shall prove by induction on the number of edges)

If $|E(G)| = |V(G)| - 1$, then G is a tree and $|F(G)| = 1$. So equality holds.

Now assume that the formula holds if $|E(G)| = |V(G)| + n - 1$ for some integer $n \geq 0$. Then consider a graph G such that $|E(G)| = |V(G)| + n$.

We know that G cannot be tree and so G must contain a cycle C . If $e \in E(C)$, then $G - \{e\}$ is still connected. So we can now apply our inductive hypothesis to say that $|V(G - \{e\})| - |E(G - \{e\})| + |F(G - \{e\})| = 2$.

Note that $F(G - \{e\}) = F(G) - 1$ because by removing e , we removed the boundary seperating two different faces, thus merging them. Additionally, $V(G) = V(G - \{e\})$ and $E(G) = E(G - \{e\}) + 1$. So we can thus conclude that: $|V(G)| - (|E(G)| - 1) + (|F(G) - 1) = 2$. Or in other words:

$$|V(G)| - |E(G)| + |F(G)| = 2$$

Theorem: Let G be a connected planar graph containing at least one cycle but containing no cycles of length less than g . Then:

$$|E(G)| \leq \frac{g}{g-2} (|V(G)| - 2)$$

Proof: We know that the degree of each face $f \in F(G)$ is at least g . Thus, by the handshaking lemma, we get that $g|F(G)| \leq 2|E(G)|$. Now, inserting this into Euler's formula, we get that:

$$|V(G)| - |E(G)| + \frac{2}{g}|E(G)| \geq |V(G)| - |E(G)| + |F(G)| = 2$$

This can then be manipulated to get the above formula.

This upper bound is maximized when $g = 3$, thus giving a an upper bound not dependent on g of:

$$|E(G)| \leq 3|V(G)| - 6$$

All in all, this is a neat necessary condition for a graph G to be planar. Beware though that this is not a sufficient condition. Some nonplanar graphs satisfy this inequality.

Some example calculations:

- K_5 : The minimum cycle in K_5 is of length 3. So note:

$$\frac{3}{3-2}(|V(G)| - 2) = 9 < 20 = |E(K_5)|$$

Thus, we have shown that K_5 can't be planar.

- Petersen Graph: The minimum cycle length in the Petersen graph is 5. So note:

$$\frac{5}{5-2}(|V(G)| - 2) = 13\frac{1}{3} < 15 = |E(K_5)|$$

Thus, we have shown that the Petersen graph can't be planar.

Lecture 9: 2/8/2024

A proper (vertex) coloring of G is a map $c : V(G) \rightarrow S$ where S is a set such that $c(u) \neq c(v)$ if $\{u, v\} \in E(G)$.

The chromatic number $\chi(G)$ is the minimum k for which G has a proper k -coloring.

A graph G is d -degenerate if every induced subgraph of G has a vertex of degree at most d .

Lemma: If G is d -degenerate, then $\chi(G) \leq d + 1$.

Proof:

Firstly, observe that if a graph G is d -degenerate, then any subgraph resulting by removing vertices from G will also be d -degenerate.

Also, observe that for any induced subgraph H of G , we clearly have that $\delta(H) \leq d$. Therefore, assuming H is nonempty, we know there exists a vertex $v \in V(H)$ such that $d(v) \leq d$.

Using the above two observations, we can do induction over a d -degenerate graph as follows:

For each integer $i \in \{1, \dots, |V(G)|\}$, remove a vertex v_i from $H_i = G - \{v_1, \dots, v_{i-1}\}$ satisfying the property that $d_{H_i}(v_i) \leq d$.

Doing this, we will eventually be left with an empty graph. This graph obviously has a proper coloring using at most $d + 1$ colors.

Next, consider adding vertices back in. For any $i \in \{1, \dots, |V(G)|\}$, we shall inductively assume that $G - \{v_1, \dots, v_{i-1}, v_i\}$ has a proper coloring using less than $d + 1$ colors. Then, because we specified that v_i has at most a degree of d in $G - \{v_1, \dots, v_{i-1}\}$, we know there is at least one available color to assign v_i . So $G - \{v_1, \dots, v_{i-1}\}$ has a proper coloring using at most $d + 1$ colors. ■

Brook's Theorem: If G is a connected graph, then $\chi(G) \leq \Delta G$ unless G is an odd cycle or a complete graph.

Proof: (we shall proceed by induction on the number of vertices in G)

To start, consider that $\chi(K_n) = \Delta(K_n) + 1$. Similarly for an odd cycle C , we have that $\chi(C) = \Delta(C) + 1$. Hence, this is why Brook's theorem specifies that G is not one of the above graphs. Additionally, note that if $\Delta G < 3$, then the theorem is obviously true. So, we only need to focus on when $\Delta G \geq 3$.

Now, note that we may assume for all $v \in V(G)$ that $G - \{v\}$ is connected.

After all, if $G - \{v\}$ was not connected, then we could write G as the union of two smaller connected graphs: G_1 and G_2 such that $V(G_1) \cap V(G_2) = \{v\}$. Now if G_i is an odd cycle or a complete graph, then we know $\chi(G_i) = \Delta(G_i) + 1$. However, then $\Delta(G_i) \leq \Delta(G) - 1$ because the degree of the vertex connecting G_i to the rest of G is equal to $\Delta(G_i)$. Meanwhile if G_i is not an odd cycle or complete graph, then we can inductively prove Brook's Theorem for G_i to get that $\chi(G_i) \leq \Delta(G_i) \leq \Delta(G)$. In either case, we will get a coloring over G_i using at most ΔG colors.

So G_1 and G_2 both have colorings with at most $\Delta(G)$ colors. By setting v in both graphs to have the same color, we can then merge the colorings of G_1 and G_2 to get a suitable coloring of G .

Thus, assume that $G - \{v\}$ is connected for all $v \in V(G)$. Then, if $G \neq K_{\Delta(G)+1}$, we can pick vertices v_1, v_{n-1} and v_n such that $\{v_1, v_{n-1}\} \in E(G)$, $\{v_1, v_n\} \in E(G)$, and $\{v_{n-1}, v_n\} \notin E(G)$.

Case 1: Suppose $H = G - \{v_n\} - \{v_{n-1}\}$ is connected. Then, we can order the vertices of H : v_1, v_2, \dots, v_{n-2} so that for $i \geq 2$, v_i always has at least one neighbor v_j where $j < i$.

This was proven in a homework exercise...

Assign v_n and v_{n-1} the color 1. Then, considering our graph G , since v_i has at most $\Delta(G) - 1$ neighbors v_j where $i < j$, we can iterate from v_{n-2} to v_2 , assigning whatever color is available to each vertex. After all, it is guaranteed that we will not have assigned all $\Delta(G)$ colors to the neighborhood of v_i yet since we could have only assigned $\Delta(G) - 1$ colors at most.

If v_i is neighbors with v_n or v_{n+1} then that just means it isn't neighbors with one or two more v_j in our ordering. So, this doesn't affect our conclusion.

Also, the number of colors assigned in $N(v_1)$ is at most $\Delta(G) - 1$ since two of v_1 's neighbors are assigned the same color. So we can assign it a remaining color afterwards. Hence G has a proper coloring.

Case 2: Now suppose $H = G - \{v_n\} - \{v_{n+1}\}$ is disconnected. Then we may view G as the union of two smaller connected graphs G_1 and G_2 such that $V(G_1) \cap V(G_2) = \{v_n, v_{n+1}\}$. By the same type of induction that we did earlier, we can prove that G_1 and G_2 can be covered by at most $\Delta(G)$ colors. So, choose a coloring such that v_n and v_{n+1} are assigned the same colors in both G_1 and G_2 . Then, we can merge the two colorings to get a proper coloring of G . ■

Note that $\Delta(G)$ can differ greatly from $\chi(G)$. For instance, bipartite graphs are always 2-colorable no matter how many edges they have.

A maximal planar graph is a graph that is planar but the addition of any edge results in a non-planar graph. Similarly, a maximal plane graph is a plane drawing of a maximal planar graph.

Theorem: Every planar graph has a vertex v of degree at most 5.

Proof:


If G is planar, then $|E(G)| \leq 3|V(G)| - 6$. However,

$$\sum_{v \in V(G)} d_G(v) = 2|E(G)| \leq 6|V(G)| - 12 < 6|V(G)|.$$

So not every vertex in G can have a degree of 6 or more.

A contraction of an edge $\{a, b\}$ in a graph G , notated as $G/\{a, b\}$, is the graph obtained by adding a vertex x to $G - \{a\} - \{b\}$ so that x is adjacent to all vertices in $N_G(a) \cup N_G(b)$. An important property of contractions is that if G is planar, then $G/\{a, b\}$ is also planar.

What a contraction looks like:



The 5-color Theorem: If G is a planar graph, then $\chi(G) \leq 5$.

Proof:

If $|V(G)| \leq 5$, then we may simply assign all vertices different colors to show that $\chi(G) \leq 5$. So, consider when $|V(G)| \geq 6$.

Then if $\delta(G) \leq 4$, we know there exists $v \in V(G)$ such that $d_G(v) \leq 4$. So, by induction we can find a proper 5-coloring of $G - \{v\}$. And because v has at most 4 neighbors, there is at least one color left over to assign v afterwards. Meanwhile, because of the previous theorem, we know that any induced subgraph of G will have a minimum degree of at most 5. So, the only non-trivial case we need to prove is when $\delta(G) = 5$.

Thus, assume there exists a vertex $v \in V(G)$ such that $d(v) = 5$. Since K_5 is not planar, there exists some pair of vertices: $a, b \in N_G(v)$ which are not adjacent. Define $H = G / \{a, b\}$ with a new vertex w adjacent to $N_G(v) \cup N_G(a)$.

Next contract H again, defining $I = H / \{b, w\}$ with a new vertex x adjacent to $N_H(w) \cup N_H(b)$. Note that I and H are planar. Thus, since I is a planar graph with less vertices than G , we can do induction to find a proper 5-coloring of I . We'll call this coloring c .


For each $y \in V(I) \cap V(G)$, define $c'(y) = c(y)$. This effectively gives a proper coloring of $G - a - b - v$. Then assign $c'(a) = c'(b) = c(x)$. We can do this because we chose a and b to not be adjacent and because none of a and b 's neighbors could have been assigned $c(x)$. And finally, because a and b in $N(v)$ have been assigned same color, there is at least one color which we can assign to v . ■

Lecture 10: 2/13/2024


The Art Gallery Problem:

Let R be a region in the plane bounded by a polygon. Also, define that two points are "*mutually visible*" if there exists a straight line between them that is contained entirely in R . Then, what is the smallest size of a set of points S in R such that every point in R is mutually visible to a point in S .


In terms of the art gallery analogy, R is the floor plan of an art gallery and $|S|$ is the number of guards you've hired to guard the gallery. Here are some example galleries with guards:




3 edges, 1 guard



4 edges, 1 guard



5 edges, 1 guard



6 edges, 2 guard

Art Gallery Theorem: For an n -sided polygon, one needs to include at most $\lfloor \frac{n}{3} \rfloor$ many points in S in order to have that every point in R be mutually visible with a point in S .

Proof:

Consider triangulating R . In other words, add straight edges between vertices of R in order to partition R into a bunch of triangular regions.

Having done that, now consider the graph G whose vertices are the vertices of R and whose edges are the boundaries of the triangles of our triangulation of R . We can prove that $\chi(G) = 3$.

When $n = 3$, we have that G consists of a single triangle. So $\chi(G) = 3$.

Now we proceed by induction. Assume that for $n > 3$, we know that $\chi(G) = 3$ if R has less than n sides. Therefore, we now consider a graph G made from an n -sided polygon. Let $\{u, v\}$ be any edge of G not on the boundary of the infinite face. Importantly, we know such a side will exist because $n > 3$. Then, we can partition G into two subgraphs G_1 and G_2 such that $V(G_1) \cap V(G_2) = \{u, v\}$ and $E(G_1) \cup E(G_2) = E(G)$.

G_1 and G_2 are the triangulations of some polygon with less than n sides. Therefore by strong induction, there exists the proper vertex 3-colorings: $c_1 : V(G_1) \rightarrow \{1, 2, 3\}$, and $c_2 : V(G_2) \rightarrow \{1, 2, 3\}$. Shifting around the colors so that $c_1(u) = c_2(u)$ and $c_1(v) = c_2(v)$, we can then define a proper 3-coloring of all G . So $\chi(G) \leq 3$.

Since G obviously can't be colored by 2 or less colors since the triangles in G need three colors, we thus conclude that $\chi(G) = 3$.

Having shown that G has a 3-coloring, now note that in that 3-coloring, the three vertices enclosing any triangular face of G must all have a different color assigned to them. Additionally, every point in a triangle is mutually visible to another point in that triangle.

Letting the least assigned color be called i , make S the set of locations of vertices which were assigned the color i . Then any point in any of the triangles partitioning R is mutually visible to a point in S . Also, $|S|$ is at most $\lfloor \frac{n}{3} \rfloor$. ■

Some other problems related to the art gallery problem are:

- The rectilinear art gallery problem.


In this version of the problem, all angles between sides must be right angles. As it turns out, with this restriction on R we only need $\lfloor \frac{n}{4} \rfloor$ points in S so that every point in R is mutually visible with a point in S . However, we aren't proving that in this class.

- The 3d art gallery problem
- The non-polygon art gallery problem

The 4-color Theorem: If G is a planar graph, then $\chi(G) \leq 4$.

Unfortunately, no short proof of this is known. The quickest proof considers 633 different cases.

Let G be a plane graph. To get the dual of G , denoted G^* , represent each face $f \in F(G)$ as a vertex in G^* and draw it inside its corresponding face in G . Then, for each edge $e \in E(G)$, draw an edge $\{f_1, f_2\} \in E(G^*)$ such that e is between the faces f_1 and f_2 of G and $\{f_1, f_2\}$ is drawn to only intercept e at one point.



Proposition: If G is connected, then $(G^*)^* \cong G$. In other words, while $(G^*)^*$ might be drawn in the plane differently from G , both graphs will have the same vertex, edge, and face sets.

The proof for this is beyond the scope of this class (and requires math I don't know yet).

The significance of the above proposition is that by taking the dual of a connected plane graph G , we can convert the faces of G into vertices and the vertices of G into faces.

Application: coloring the faces of G is equivalent to coloring the vertices of G^* .

Here is one notable approach to the four color theorem:

Firstly, to prove the 4-color theorem, it is enough to prove it for maximal planar graphs (i.e. a planar graph with as many edges as possible). After all, removing edges does not make a coloring invalid.

Also, note that in a maximal planar graph G with at least three vertices, every face (including the infinite face) is a triangle. Therefore, G^* is 3-regular. Furthermore, if G has at least 4 vertices, then no face can share more than one edge with any other face. So G^* is a simple (a.k.a non-pseudo) graph.

A graph G is k -connected if for every set X of less than k vertices, $G - X$ is still connected. Note that, if a maximal planar graph G has more than 3 vertices, then G^* is 3-connected.

Here's a quick useful calculation for the homework that the professor gave us:

Let G be a maximal planar graph with n vertices. By the handshaking lemma, we know that $3|F(G)| = 2|E(G)|$. And since G is obviously connected, we can thus say that: $|V(G)| - |E(G)| + |F(G)| = n - \frac{1}{2}|F(G)| = 2$. So, $|F(G)| = 2n - 4$. Plugging this into Euler's formula again, we then get that $|E(G)| = 3n - 6$.

As a result, we know that G^* has $2n - 4$ vertices and $3n - 6$ edges.

Lecture 11: 2/15/2024

Let \mathcal{F} be any family of graphs.

The external numbers / Turán numbers for \mathcal{F} are the quantities $\text{ex}(n, \mathcal{F})$ which denote the maximum number of edges in an n vertex graph not containing any graph in \mathcal{F} .

We call a graph that does not contain any member of \mathcal{F} an \mathcal{F} -free graph.

A k -core of a graph G is the graph obtained through the following algorithm:

1. Let X be the set of vertices removed from G in previous steps.
2. If there is no vertex $v \in V(G - X)$ with degree less than k in $G - X$, then you are done as $G - X$ is a k -core of G .
3. Otherwise, add v to the set X and go back to step 2.

If not empty, a k -core is the largest subgraph of G with minimum degree at least k .

The order in which vertices are removed does not matter. The k -core of a graph is unique.

Lemma: If G is an n -vertex graph with more than $(k-1)n - \binom{k}{2}$ edges, then G has a nonempty k -core.

Proof:

Now, if $\delta(G) \geq k$, we're done. So, assume that there is a vertex $v_1 \in V(G)$ with $d(v_1) \leq k-1$. Then:

$$|E(G - \{v\})| > (k-1)n - \binom{k}{2} - (k-1) = (k-1)(n-1) - \binom{k}{2}$$

Now if $n = k$, then $|E(G)| > (k-1)k - \binom{k}{2} = \binom{k}{2}$. But that would mean that G has more edges than K_n , which is impossible. So, we know that $n > k$. Therefore, it makes sense to consider repeating the above step $n - k$ times. Having done that, we will get that there are more than $(k-1)(k) - \binom{k}{2}$ edges remaining. However, that is impossible as k vertex graph can't have more than $\binom{k}{2}$ edges. Hence, we conclude that the algorithm must have terminated before all $n - k$ steps could be done.

Corollary: If G is a graph with n vertices and more than $(k-1)n - \binom{k}{2}$ edges, then G contains a cycle of length at least $k+1$ unless $G = K_k$.

Proof:

Let $H \subseteq G$ be a non-empty k -core in G . Since $\delta(H) \geq k$, H contains a cycle of length at least $k+1$. Thus so does G . ■

A cut: (A, B) , of a graph G is a spanning subgraph of G such that $V(G)$ is partitioned into two sets A and B and an edge $\{a, b\} \in E(G)$ is only included in the cut if $a \in A$ and $b \in B$.

A max cut is a cut with as many edges as possible.

Theorem: Every graph with m edges has a cut with at least $m/2$ edges.

Proof:

Let G have m edges. On each vertex of G , flip a fair coin. Then, let A be the set of vertices which have heads and let B be the set of vertices which have tails. That way, (A, B) is a cut.

Since the assignment of each vertex is independent from the assignment of any other vertex, we have that there is a $\frac{1}{2}$ probability of any edge being included in the cut. Thus, the mean number of edges between A and B is:

$$\sum_{e \in E(G)} \mathbb{P}(e \text{ is in the cut}) = \sum_{i=1}^m \frac{1}{2} = \frac{1}{2}m$$

This implies that some cut (A, B) of G must have at least $m/2$ edges. (It also implies some cut has at most $m/2$ edges.)

Erdős-Gallai Theorem: Let $k \geq 1$, $n \geq 1$, and let G be an n -vertex P_k -free graph. Then $|E(G)| \leq (k-1)\frac{n}{2}$ with equality if and only if k divides n and every component of G is K_k .

Proof:

Let G be an n -vertex graph with at least $(k-1)\frac{n}{2}$ edges.

Well frick. Once again I do not understand how his proof works. Why does every component need to be K_k ?!